

# Progetto - Deep Learning on Temporal Data

Confronto ARIMA model e ARMA(1,1)-sGARCH(1,1)

Davide Zicca

9/7/2021

## Motivazione e data setting

L'analisi che viene qui proposta prende spunto dalle lezioni *Deep Learning on Temporal Data* e presenta un confronto tra i modelli ARIMA e un modello ARMA(1,1)-sGARCH(1,1) che presenta dunque anche l'equazione per la varianza. I modelli GARCH sono largamente usati in letteratura per catturare gli effetti *disruptive* a cui sono spesso soggetti i ticker azionari. Ho scelto i prezzi azionari di Amazon da Giugno 2015 a Giugno 2021.

```
#librerie necessarie
library(ggplot2)
library(tidyquant)
library(timetk)
library(tseries)
library(timeSeries)
library(forecast)
library(seastests)
library(rugarch)
library(fDMA)
library(dplyr)
set.seed(29)
amazon = tq_get("AMZN",
                from = '2015-06-01',
                to = "2021-06-01",
                get = "stock.prices")
```

## Data Visualization, calcolo log returns e fit del MACD

```
amazon %>%
  ggplot(aes(x = date, y = adjusted)) +
  geom_line() +
  ggtitle("Amazon") +
  labs(x = "Date", "Price") +
  scale_x_date(date_breaks = "years", date_labels = "%Y") +
  labs(x = "Data", y = "Adjusted Price") +
  theme_minimal()
```




```
AMZN_annual_returns = amazon %>%
  group_by(symbol) %>%
  tq_transmute(select      = adjusted,
                mutate_fun = periodReturn,
                period      = "yearly",
                type         = "arithmetic")
AMZN_annual_returns
```

```
## # A tibble: 7 x 3
## # Groups:   symbol [1]
##   symbol date      yearly.returns
##   <chr> <date>         <dbl>
## 1 AMZN  2015-12-31      0.568
## 2 AMZN  2016-12-30      0.109
## 3 AMZN  2017-12-29      0.560
## 4 AMZN  2018-12-31      0.284
## 5 AMZN  2019-12-31      0.230
## 6 AMZN  2020-12-31      0.763
## 7 AMZN  2021-05-28     -0.0104
```

```
AMZN_annual_returns %>%
  ggplot(aes(x = date, y = yearly.returns, fill = symbol)) +
  geom_col() +
  geom_hline(yintercept = 0, color = palette_light()[[1]]) +
  scale_y_continuous(labels = scales::percent) +
  labs(title = "Amazon: Annual Returns",
       y = "Annual Returns", x = "") +
```


```
facet_wrap(~ symbol, ncol = 2, scales = "free_y") +
theme_minimal() +
scale_fill_tq()
```



```
AMZN_daily_log_returns = amazon %>%
  group_by(symbol) %>%
  tq_transmute(select      = adjusted,
                mutate_fun = periodReturn,
                period      = "daily",
                type         = "log",
                col_rename   = "daily.returns")

AMZN_daily_log_returns %>%
  ggplot(aes(x = daily.returns, fill = symbol)) +
  geom_density(alpha = 0.5) +
  labs(title = "Amazon: Charting the Daily Log Returns",
       x = "Daily Returns", y = "Density") +
  theme_minimal() +
  scale_fill_tq() +
  facet_wrap(~ symbol, ncol = 2)
```

## Amazon: Charting the Daily Log Returns



Calcolo e plot del Moving Average Convergence Divergence:


```
AMZN_macd = amazon %>%
  group_by(symbol) %>%
  tq_mutate(select      = close,
            mutate_fun = MACD,
            nFast      = 12,
            nSlow      = 26,
            nSig       = 9,
            maType     = SMA) %>%
  mutate(diff = macd - signal) %>%
  select(-(open:volume))
AMZN_macd
```

```
## # A tibble: 1,511 x 6
## # Groups:   symbol [1]
##   symbol date      adjusted macd signal diff
##   <chr> <date>      <dbl> <dbl> <dbl> <dbl>
## 1 AMZN 2015-06-01    431.    NA    NA    NA
## 2 AMZN 2015-06-02    431.    NA    NA    NA
## 3 AMZN 2015-06-03    437.    NA    NA    NA
## 4 AMZN 2015-06-04    431.    NA    NA    NA
## 5 AMZN 2015-06-05    427.    NA    NA    NA
## 6 AMZN 2015-06-08    424.    NA    NA    NA
## 7 AMZN 2015-06-09    425.    NA    NA    NA
## 8 AMZN 2015-06-10    431.    NA    NA    NA
## 9 AMZN 2015-06-11    433.    NA    NA    NA
```

```
## 10 AMZN    2015-06-12    430.    NA    NA    NA
## # ... with 1,501 more rows
```

```
AMZN_macd %>%
  filter(date >= as_date("2021-05-01")) %>%
  ggplot(aes(x = date)) +
  geom_hline(yintercept = 0, color = palette_light()[[1]]) +
  geom_line(aes(y = macd, col = symbol)) +
  geom_line(aes(y = signal, color = "blue", linetype = 2) +
  geom_bar(aes(y = diff), stat = "identity", color = palette_light()[[1]]) +
  facet_wrap(~ symbol, ncol = 2, scale = "free_y") +
  labs(title = "AMZN: Moving Average Convergence Divergence",
       y = "MACD", x = "", color = "") +
  theme_minimal() +
  scale_color_tq()
```

## AMZN: Moving Average Convergence Divergence



## Test: stazionarietà e presenza ARCH effect

```
#Stationarity test
#adf.test(AMZN_daily_log_returns$daily.returns, alternative = "stationary")
# p-value = 0.01 -> si rifiuta l'ipotesi nulla
ts= ts(AMZN_daily_log_returns)[,3]

# seasonality test
isSeasonal(ts, freq=1)
```

```
## [1] FALSE
#stationary test
adf.test(ts, alternative= "stationary")

##
## Augmented Dickey-Fuller Test
##
## data:  ts
## Dickey-Fuller = -11.801, Lag order = 11, p-value = 0.01
## alternative hypothesis: stationary
#ARCH EFFECT

archtest(as.vector(AMZN_daily_log_returns$daily.returns))

##
## Engle's LM ARCH Test
##
## data:  as.vector(AMZN_daily_log_returns$daily.returns)
## statistic = 53.451, lag = 1, p-value = 2.652e-13
## alternative hypothesis: ARCH effects of order 1 are present
# presenza ARCH effect
```

## Fit modelli ARMA, verifica auto-correlation e valori BIC per scelta modello

```
n = length(ts)
n

## [1] 1511
nV=round(n/3) # Validation set (33% del totale)
nV

## [1] 504
nT=n-nV # training set - osservazioni
train=ts[c(1:nT)]
valid=ts[c((nT+1):n)]

auto_model1=auto.arima(train, ic="aic", stationary=FALSE,seasonal=FALSE)
auto_model1

## Series: train
## ARIMA(0,0,0) with non-zero mean
##
## Coefficients:
##      mean
##      0.0014
## s.e.  0.0006
##
## sigma^2 estimated as 0.0003423:  log likelihood=2589.54
## AIC=-5175.08   AICc=-5175.07   BIC=-5165.25
```

```
res1=auto_model1$residuals
```

```
Box.test(res1, lag = 12, type = "Box-Pierce", fitdf=0)
```

```
##
```

```
## Box-Pierce test
```

```
##
```

```
## data: res1
```

```
## X-squared = 12.998, df = 12, p-value = 0.3692
```

```
Box.test(res1, lag = 12, type = "Ljung-Box", fitdf=0)
```

```
##
```

```
## Box-Ljung test
```




```
##
```

```
## data: res1
```

```
## X-squared = 13.104, df = 12, p-value = 0.3616
```

```
checkresiduals(auto_model1, include.mean=FALSE, plot=TRUE)
```

### Residuals from ARIMA(0,0,0) with non-zero mean



```
##
```

```
## Ljung-Box test
```

```
##
```


```
## data: Residuals from ARIMA(0,0,0) with non-zero mean
```

```
## Q* = 12.87, df = 9, p-value = 0.1686
```

```
##
```

```
## Model df: 1. Total lags used: 10
```


```
acf(train, lag=10)
```



```
pacf(train, lag=10)
```



## Series train



```
arima_model=arima(train, order=c(1,0,1) )
arima_model
```

```
##
## Call:
## arima(x = train, order = c(1, 0, 1))
##
## Coefficients:
##      ar1      ma1  intercept
##      0.0051  0.0062    0.0014
## s.e.  1.2736  1.3051    0.0006
##
## sigma^2 estimated as 0.0003419:  log likelihood = 2589.6,  aic = -5171.2
```

```
BIC(auto_model1)
```

```
## [1] -5165.248
```

```
BIC(arima_model)
```

```
## [1] -5151.545
```

## Modello ARMA(1,1)-sGARCH(1,1): fit, risultati e plot

```
# GARCH MODEL assumendo una distribuzione Normale
#sGARCH
library(rugarch)
```

```

s_garchMod = ugarchspec(mean.model = list(armaOrder = c(1, 1), include.mean = TRUE
),
variance.model = list(model = 'sGARCH',
                      garchOrder = c(1, 1)),
distribution.model = "norm")

s_garchFit = ugarchfit(spec=s_garchMod, data=ts)
s_garchFit

```

```

##
## *-----*
## *          GARCH Model Fit          *
## *-----*
##
## Conditional Variance Dynamics
## -----
## GARCH Model   : sGARCH(1,1)
## Mean Model    : ARFIMA(1,0,1)
## Distribution   : norm
##
## Optimal Parameters
## -----
##      Estimate  Std. Error  t value Pr(>|t|)
## mu      0.002084    0.000315   6.6101     0
## ar1     0.797653    0.114210   6.9841     0
## ma1    -0.840888    0.101013  -8.3246     0
## omega   0.000027    0.000005   5.3200     0
## alpha1  0.181082    0.028111   6.4417     0
## beta1   0.754320    0.030798  24.4923     0
##
## Robust Standard Errors:
##      Estimate  Std. Error  t value Pr(>|t|)
## mu      0.002084    0.000405   5.1413 0.00000
## ar1     0.797653    0.127935   6.2348 0.00000
## ma1    -0.840888    0.111760  -7.5241 0.00000
## omega   0.000027    0.000011   2.5527 0.01069
## alpha1  0.181082    0.035858   5.0500 0.00000
## beta1   0.754320    0.043618  17.2939 0.00000
##
## LogLikelihood : 3999.049
##
## Information Criteria
## -----
##
## Akaike          -5.2853
## Bayes           -5.2642
## Shibata         -5.2853
## Hannan-Quinn    -5.2774
##
## Weighted Ljung-Box Test on Standardized Residuals
## -----
##
##              statistic  p-value
## Lag[1]              3.267 0.070701

```

```

## Lag[2*(p+q)+(p+q)-1][5]      4.934 0.004145
## Lag[4*(p+q)+(p+q)-1][9]      5.617 0.324344
## d.o.f=2
## H0 : No serial correlation
##
## Weighted Ljung-Box Test on Standardized Squared Residuals
## -----
##
##                statistic p-value
## Lag[1]                0.3363 0.5620
## Lag[2*(p+q)+(p+q)-1][5]    0.6483 0.9326
## Lag[4*(p+q)+(p+q)-1][9]    1.0435 0.9844
## d.o.f=2
##
## Weighted ARCH LM Tests
## -----
##
##      Statistic Shape Scale P-Value
## ARCH Lag[3]    0.06644 0.500 2.000 0.7966
## ARCH Lag[5]    0.30912 1.440 1.667 0.9378
## ARCH Lag[7]    0.44368 2.315 1.543 0.9832
##
## Nyblom stability test
## -----
## Joint Statistic: 0.6034
## Individual Statistics:
## mu      0.14755
## ar1      0.03782
## ma1      0.04302
## omega    0.05536
## alpha1   0.05712
## beta1    0.06437
##
## Asymptotic Critical Values (10% 5% 1%)
## Joint Statistic:      1.49 1.68 2.12
## Individual Statistic: 0.35 0.47 0.75
##
## Sign Bias Test
## -----
##
##                t-value   prob sig
## Sign Bias      1.992939 0.04645  **
## Negative Sign Bias 0.594396 0.55234
## Positive Sign Bias 0.000715 0.99943
## Joint Effect    5.879244 0.11764
##
##
## Adjusted Pearson Goodness-of-Fit Test:
## -----
##      group statistic p-value(g-1)
## 1      20      90.5    2.709e-11
## 2      30     104.9    1.567e-10
## 3      40     109.5    1.313e-08
## 4      50     121.3    4.648e-08
##
##
## Elapsed time : 0.433579

```


```
## Risultati del modello
```

```
coef(s_garchFit)
```

```
##          mu          ar1          ma1          omega          alpha1  
## 2.083982e-03 7.976531e-01 -8.408882e-01 2.717585e-05 1.810823e-01  
##          beta1  
## 7.543195e-01
```


```
s_rhat = s_garchFit@fit$fitted.values
```

```
plot.ts(s_rhat)
```



```
s_hhat = ts(s_garchFit@fit$sigma^2)
```

```
plot.ts(s_hhat)
```



```

fit.val      = coef(s_garchFit)
fit.sd       = diag(vcov(s_garchFit))
true.val     = s_garchFit@fit$tval

fit.conf.lb  = fit.val + qnorm(0.025) * fit.sd
fit.conf.ub  = fit.val + qnorm(0.975) * fit.sd
print(fit.val)

##          mu          ar1          ma1          omega          alpha1
## 2.083982e-03  7.976531e-01 -8.408882e-01  2.717585e-05  1.810823e-01
##          beta1
## 7.543195e-01

print(fit.sd)

## [1] 9.939810e-08 1.304401e-02 1.020362e-02 2.609369e-11 7.902315e-04
## [6] 9.485306e-04

print(true.val)

##          mu          ar1          ma1          omega          alpha1          beta1
## 6.610052  6.984067 -8.324558  5.320047  6.441676 24.492306

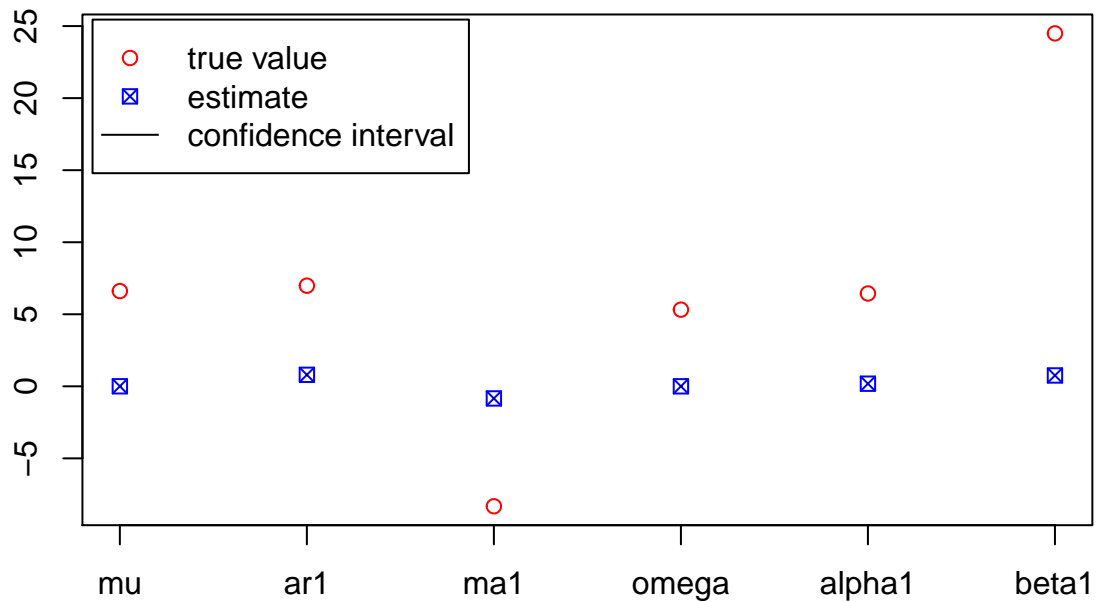
plot(true.val, pch = 1, col = "red",
      ylim = range(c(fit.conf.lb, fit.conf.ub, true.val)),
      xlab = "", ylab = "", axes = FALSE)
box(); axis(1, at = 1:length(fit.val), labels = names(fit.val)); axis(2)
points(coef(s_garchFit), col = "blue", pch = 7)

```

```

for (i in 1:length(fit.val)) {
  lines(c(i,i), c(fit.conf.lb[i], fit.conf.ub[i]))
}
legend( "topleft", legend = c("true value", "estimate", "confidence interval"),
       col = c("red", "blue", 1), pch = c(1, 7, NA), lty = c(NA, NA, 1), inset = 0.01)

```



```

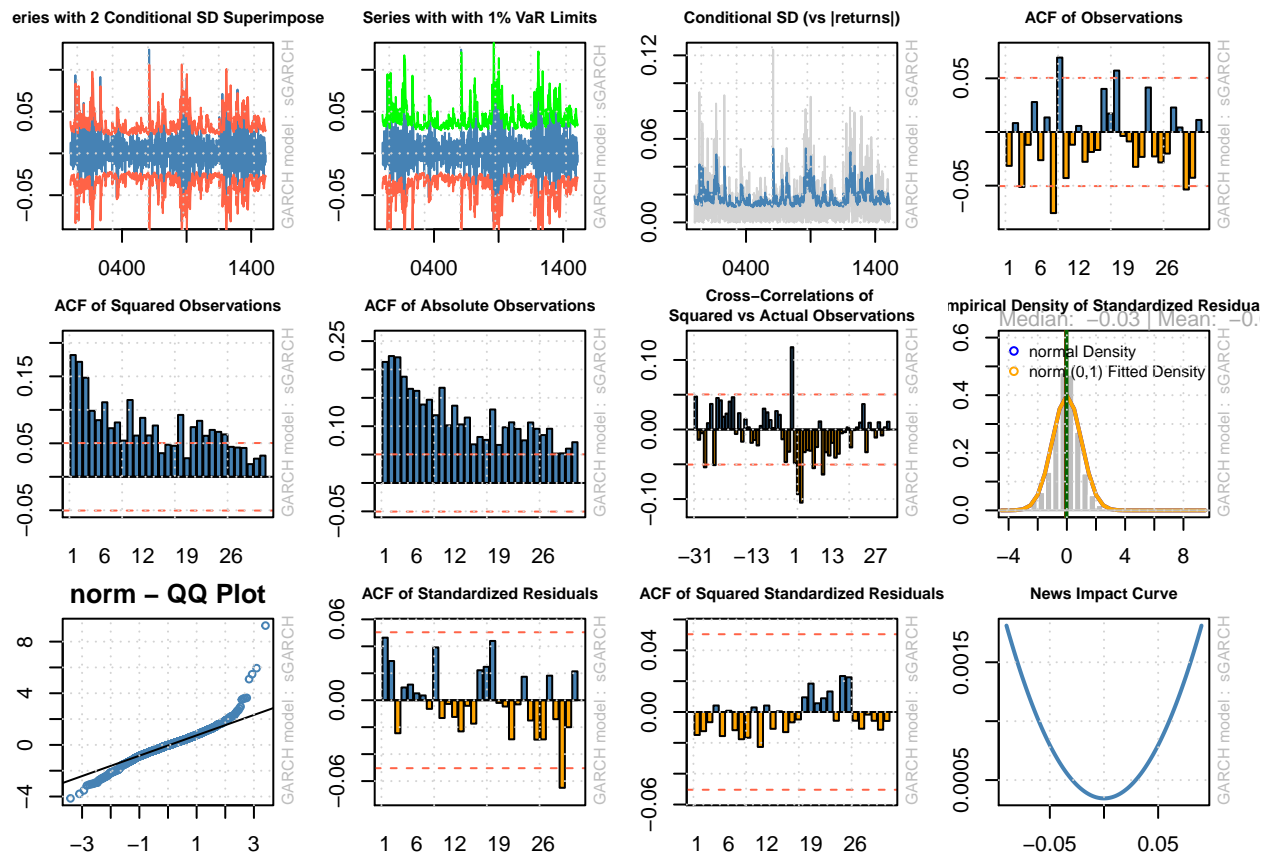
par(mfrow=c(2, 3))
par(mar = c(2, 2, 2, 2))
plot(s_garchFit, which="all")

```

```

##
## please wait...calculating quantiles...

```



## Confronto tra modelli

```
# confronto tra modelli
infocriterio(s_garchFit)[2] #BIC= -5.264179

## [1] -5.264179

BIC(auto_model1) # = -5165.248 -> è il valore più basso e quindi quello da

## [1] -5165.248

# scegliere. auto_model1 è il modello da scegliere
BIC(arima_model) # = -5151.545

## [1] -5151.545
```

## Conclusione

Dopo una attenta analisi dei valori BIC, viene scelto il modello `auto_model1` poichè mostra il valore più basso di BIC. Tuttavia, il modello  $\text{ARMA}(1,1)\text{-sGARCH}(1,1)$ , grazie alla presenza dell'equazione della varianza è in grado di fornire informazioni maggiori rispetto ai modelli ARMA che presentano la sola equazione della media.